

Parallel Distributed Processing in 21st-Century Neuroscience and AI

Ishan Mysore
ATCS: Neural Networks, Period 3
December 8, 2025

1. Introduction

A single neuron in the visual cortex tells us almost nothing about what we are seeing, yet the coordinated firing of millions of such neurons generates our rich visual experience of the world. This fundamental insight—that intelligence emerges not from individual specialized components but from vast networks of simple units working in parallel—has become the organizing principle uniting 21st-century neuroscience and artificial intelligence.

This paper explores the convergence of these two fields under the parallel distributed processing (PDP) framework—a convergence made possible by several key developments in artificial intelligence: exponential increases in computational power, the availability of massive labeled datasets, and breakthroughs in optimization algorithms. On the neuroscience side, high-density recording technologies now capture thousands of neurons simultaneously in behaving animals, generating the rich population data needed for meaningful comparisons between AI models and biological systems. Together, these technological and methodological advances have rendered PDP crucial in understanding how both biological and artificial systems process information.

2. A Brief Overview of PDP

Before examining these developments in detail, it is essential to define what we mean by parallel distributed processing. In the landmark work *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, Rumelhart and McClelland formalized PDP as a computational framework in which cognition emerges from many simple processing units (neurons or artificial nodes) working in parallel [1]. In a PDP system, no single unit contains a complete representation; instead, meaning arises from the collective pattern of activation across the entire population [1]. This distributed encoding provides two key advantages: graceful degradation—when some units fail, performance declines proportionally rather than collapsing entirely—and the ability to generalize from partial or noisy inputs [2].

Additionally, the PDP framework stands in contrast to traditional symbolic AI approaches that store knowledge in discrete symbols and rules. Instead, PDP systems learn by adjusting

the strengths of connections between units, gradually shaping the response patterns of the network through experience [2].

3. Population Coding and Neural Manifolds

The most fundamental idea of the PDP framework is that complex information cannot be decoded from any single neuron but emerges from distributed patterns across many units. In neuroscience, this is called *population coding* [3]. In a study by Xu et al., when a cat localizes a sound source, ensembles of auditory cortical neurons—each broadly tuned to different spatial positions—fire in specific patterns that collectively encode the location of the sound with much greater precision than any single neuron could achieve [4]. This distributed encoding has enabled breakthrough applications: brain-computer interfaces now decode intended movements from coordinated neural activity, allowing paralyzed patients to control robotic limbs [5]. This exemplifies the core PDP principle that intelligence arises from collective patterns rather than individual specialized units.

Neuroscientists have found that although we might record from hundreds or thousands of neurons, the meaningful patterns of population activity occupy only a few dimensions—what mathematicians call a low-dimensional *neural manifold* [6] [7]. Even though there are many neurons, the system effectively uses only a few key dimensions to encode behaviorally relevant information—which direction the animal is moving, what object it's seeing, or where a sound is coming from [7]. This concept is illustrated in Figure 1.




Figure 1. Neural population activity occupies a low-dimensional manifold. Although recordings capture activity from hundreds of neurons (high-dimensional space), behaviorally relevant patterns lie on a structured low-dimensional

surface. Here, neural activity during a reaching task traces organized trajectories in a 3D projection of the full neural space. Adapted from Gosztolai et al. (2025).

Gosztolai et al. found that both biological brains and artificial neural networks independently converge on these low-dimensional representations [7]. This discovery has transformed neural engineering: understanding manifold structure guides the design of efficient artificial networks that match performance with fewer parameters [8]. The result that vastly different systems—whether biological or artificial—discover the same solution provides compelling evidence that distributed, low-dimensional coding is a fundamental principle of intelligence.

Crucially, this compression is itself a form of distributed processing: no single neuron carries complete information, yet the coordinated low-dimensional activity of the population efficiently represents complex behavioral states. This finding supports a core PDP principle—that intelligence emerges not from adding more units, but from organizing many simple units into structured, coordinated patterns.

4. Brain Imaging and Recording Advances

Testing this principle in real populations requires recording from many neurons simultaneously—a capability that has only recently become feasible. High-density electrode arrays such as Neuropixels probes now record from nearly 1,000 neurons simultaneously across multiple brain regions in behaving mice, while two-photon calcium imaging captures activity from 10,000 or more neurons in a single field of view [9]. In humans, functional MRI and intracranial electrophysiology capture whole-brain or multi-region population signals, potentially enabling biomarkers for the early detection of psychiatric or neurological disorders [10].

These technological breakthroughs have enabled transformative applications: neurosurgeons now use population recordings to map critical brain regions before tumor removal with unprecedented precision, while researchers developing brain-computer interfaces decode complex motor intentions by analyzing coordinated activity patterns that would be invisible in single-neuron recordings [9] [10]. By applying machine learning tools to high-dimensional data, scientists can uncover underlying representational structure and directly compare biological population codes with artificial neural network activations [11]. This capability to quantitatively align brain and model responses has transformed PDP into a testable framework with measurable predictions, driving both neuroscience discoveries and improvements in AI architectures.

5. Deep Learning Architectures

Having established how neuroscientists measure and characterize distributed population codes in biological brains, we now turn to the artificial side of the PDP framework: the deep learning architectures that implement similar principles in silicon.

5.1 Convolutional Networks and the Visual Brain

Convolutional neural networks (CNNs) exemplify PDP principles in modern AI vision systems. Each layer applies many parallel filters to an input image, with the collective activations forming a distributed representation [12]. As data propagate through the network, filters become tuned to features of increasing complexity—edges and textures in early layers, shapes and object parts in deeper layers—as shown in Figure 2. Remarkably, this feature hierarchy closely mirrors the primate ventral visual stream, which is responsible for identifying objects and their qualities. Lower CNN layers explain neural activity in early visual areas (V1, V2) while deeper layers correspond to higher visual regions (V4, IT cortex) [13].




Figure 2. Hierarchical correspondence between CNN layers and human visual cortex. Neural activity patterns in different visual brain areas (Brodman areas 17, 18, 19, 37, and 20, listed vertically) are compared with activity patterns in different CNN layers (layers 0–8, listed horizontally). Higher values indicate stronger correlations. The diagonal pattern in the CNN results reflects the pattern seen in biological systems: early visual areas (V1, V2) align with early CNN layers that process simple features, while higher visual regions (V4, IT cortex) align with deeper CNN layers that represent more complex information. This correspondence demonstrates that CNNs and biological vision systems independently converge on similar hierarchical processing strategies. Adapted from Kuzovkin et al. (2018).

This precise brain-CNN correspondence has driven practical breakthroughs: autonomous vehicle vision systems now use CNN architectures specifically designed to mimic the robust processing of human vision, achieving better performance in challenging conditions like poor lighting or partial occlusion [14]. Medical imaging systems inspired by this correspondence detect cancers and retinal diseases with accuracy matching or exceeding human radiologists [15]. This discovery—artificial networks independently converge on the

same representations as biological evolution—demonstrates the importance and significance of hierarchical, distributed processing in the 21st century.

5.2 Transformer Networks and Language

Transformers process language through self-attention mechanisms where each word position compares itself to all other positions via multiple parallel "attention heads" (as shown in Figure 3). This approach enables the model to encode long-range dependencies like subject-verb agreement or thematic relationships spanning paragraphs [16].




Figure 3. Transformer self-attention patterns during sentence processing. The matrix shows which words (horizontal axis) each position (vertical axis) attends to when encoding a sentence. Different attention heads capture different linguistic relationships, including syntactic dependencies (such as subject–verb agreement) and thematic roles. Higher values indicate stronger attention weights. Adapted from Kumar et al. (2024).

Kumar et al. (2024) found that different attention heads in a transformer model predicted activity in different language brain regions during story comprehension [16]. This one-to-one mapping suggests that transformers, like the human brain, decompose language processing into specialized modules—some handling syntax, others semantics—providing evidence that distributed specialization emerges naturally in both biological and artificial language systems.

This alignment between model and brain has driven transformative applications, including real-time translation systems that serve billions of users, clinical documentation tools that reduce physician burnout by generating medical notes automatically, and accessibility

technologies for nonverbal individuals [17]. Similarly, insights into transformer attention mechanisms have led to the development of interpretability tools that help researchers understand what language models have learned and identify potential biases or errors in their reasoning. For instance, attention visualization tools now allow developers to see which words a model focuses on when generating specific outputs, making these systems more transparent and debuggable [18].

5.3 Brain-Model Alignment

Researchers now directly test PDP hypotheses by comparing neural networks to brain data through methods like representational similarity analysis (RSA). In RSA, researchers show the same set of images to both the brain and model, then measure how similarly each system responds to every pair of images. This creates two "similarity maps" that can be compared: if the brain treats two images as similar and the model does too, their representational structures align [19].

Initiatives like Brain-Score compile multiple neural and behavioral benchmarks to quantify how "brain-like" different models are, and these comparisons have revealed both striking convergences and critical gaps [20]. Schrimpf et al. (2020) found that while state-of-the-art vision models capture many aspects of primate brain activity, optimizing solely for task performance does not guarantee brain-like processing, as shown in Figure 4.




Figure 4. Brain-Score measures how well a model matches brain activity and human behavior. ImageNet performance reflects accuracy at classifying images into 1,000 object categories, a standard computer vision benchmark. Each point represents a different CNN architecture. For earlier-generation models, Brain-Score strongly correlates with ImageNet performance ($r = 0.92$), suggesting that better task performance is associated with more brain-like processing. However, for state-of-the-art models, this relationship no longer holds, indicating that optimizing for task accuracy alone does not ensure brain-like computation. Several moderately performing models, such as DenseNet-169 and ResNet-101, achieve higher Brain-Scores than top ImageNet performers. Adapted from Schrimpf et al. (2020).

These mismatches have driven important AI improvements. By recognizing that networks require vastly more labeled examples than humans, researchers have spurred the development of few-shot learning methods and self-supervised approaches that better capture how biological systems learn efficiently [21]. Additionally, adversarial vulnerability research—sparked by the discovery that networks are fooled by perturbations imperceptible to humans—has produced more robust AI systems critical for deployment in security, autonomous vehicles, and healthcare diagnostics [22]. These brain-model comparisons thus create a virtuous cycle: neuroscience constraints reveal AI weaknesses, driving engineering solutions that produce both more capable and more biologically plausible systems.

Conclusion

Modern research validates PDP as an enduring framework that continues to guide advances in both neuroscience and artificial intelligence. Both neuroscientists and AI researchers treat cognition as emerging from large networks of simple units with distributed representations. Deep convolutional models trained on object recognition develop layered feature representations that closely match the primate ventral visual stream, while transformer-based language models exhibit computations aligned with the human language network. Initiatives like Brain-Score illustrate how iteratively comparing models to neural data can drive progress and reveal where models and brains differ. Together, these findings underscore that distributed architectures and learning rules inspired by the brain can accelerate AI progress, while deep networks serve as testable, falsifiable models of brain computation. In short, focusing on these distributed designs drives rapid advances in both neuroscience and machine learning.

Looking ahead, the PDP framework suggests a virtuous cycle with concrete next steps: neuroscience discoveries about how cortex learns temporal sequences may inform the next generation of recurrent architectures, while AI advances in few-shot learning could reveal computational principles the brain uses to learn from minimal examples. As both fields continue to mature, PDP provides a testable framework for understanding intelligence in all its forms.

References

- [1] Rumelhart, D. E., & McClelland, J. L. (Eds.). (1986). *Parallel distributed processing: Explorations in the microstructure of cognition* (Vols. 1-2). MIT Press.
- [2] Rogers, T. T., & McClelland, J. L. (2003). *Semantic cognition: A parallel distributed processing approach*. MIT Press. https://stanford.edu/~jlmcc/papers/RogersMcCBook_7_03.pdf
- [3] Shamir, M. (2014). Emerging principles of population coding: In search for the neural code. *Current Opinion in Neurobiology*, 25, 140–148. <https://doi.org/10.1016/j.conb.2014.01.002>
- [4] Xu, L., Furukawa, S., & Middlebrooks, J. C. (1999). Auditory cortical responses in the cat to sounds that produce spatial illusions. *Nature*, 399(6737), 688–691. <https://doi.org/10.1038/21424>
- [5] U.S. National Institutes of Health. (2024, September 10). *Brain-computer interface helps paralyzed man speak*. NIH Research Matters. <https://www.nih.gov/news-events/nih-research-matters/brain-computer-interface-helps-paralyzed-man-speak>
- [6] Mitchell-Heggs, R., Prado, S., Gava, G. P., Go, M. A., & Schultz, S. R. (2023). Neural manifold analysis of brain circuit dynamics in health and disease. *Journal of computational neuroscience*, 51(1), 1–21. <https://doi.org/10.1007/s10827-022-00839-3>
- [7] Gosztolai, A., Peach, R.L., Arnaudon, A. *et al.* MARBLE: interpretable representations of neural population dynamics using geometric deep learning. *Nat Methods* 22, 612–620 (2025). <https://doi.org/10.1038/s41592-024-02582-2>
- [8] Edelman, B. J., Zhang, S., Schalk, G., Brunner, P., Muller-Putz, G., Guan, C., & He, B. (2025). Non-Invasive Brain-Computer Interfaces: State of the Art and Trends. *IEEE reviews in biomedical engineering*, 18, 26–49. <https://doi.org/10.1109/RBME.2024.3449790>
- [9] Grienberger, C., Giovannucci, A., Zeiger, W., & Portera-Cailliau, C. (2022). Two-photon calcium imaging of neuronal activity. *Nature reviews. Methods primers*, 2(1), 67. <https://doi.org/10.1038/s43586-022-00147-1>
- [10] Huang, Y., Wei, P.H., Xu, L. *et al.* Intracranial electrophysiological and structural basis of BOLD functional connectivity in human brain white matter. *Nat Commun* 14, 3414 (2023). <https://doi.org/10.1038/s41467-023-39067-3>
- [11] Alam, M. N. U., Basava, K., Chitransh, A., Fattah, H. M. A., Garcia-Verdugo, H. D., Lo, S. H., Lohchab, T., Martinet, K. M., Román-Palacios, C., Salazar, J. C., & Van Boxel, D. (2025). Machine learning in biological research: key algorithms, applications, and future directions. *BMC biology*, 23(1), 324. <https://doi.org/10.1186/s12915-025-02424-3>

- [12] Bowers, J. S. (2017). Parallel distributed processing theory in the age of deep networks. *Trends in Cognitive Sciences*, 21(12), 950–961. <https://doi.org/10.1016/j.tics.2017.09.013>
- [13] Kuzovkin, I., Vicente, R., Petton, M. *et al.* Activations of deep convolutional neural networks are aligned with gamma band activity of human visual cortex. *Commun Biol* 1, 107 (2018). <https://doi.org/10.1038/s42003-018-0110-y>
- [14] Abiyev, R., & Arslan, M. (2023). Vehicle detection systems for intelligent driving using deep convolutional neural networks. *Discover artificial intelligence*, 3(1), 16. <https://doi.org/10.1007/s44163-023-00062-8>
- [15] Pinto-Coelho, L. (2023). How Artificial Intelligence Is Shaping Medical Imaging Technology: A Survey of Innovations and Applications. *Bioengineering*, 10(12), 1435. <https://doi.org/10.3390/bioengineering10121435>
- [16] Kumar, S., Sumers, T.R., Yamakoshi, T. *et al.* Shared functional specialization in transformer-based language models and the human brain. *Nat Commun* 15, 5523 (2024). <https://doi.org/10.1038/s41467-024-49173-5>
- [17] Samraik, M. (2025, October 17). *AI scribes reduce physician burnout and return focus to the patient*. Yale School of Medicine. <https://medicine.yale.edu/news-article/ai-scribes-reduce-physician-burnout-return-focus-to-the-patient/>
- [18] Yeh, C., Chen, Y., Wu, A., Chen, C., Viégas, F., & Wattenberg, M. (2023). *AttentionViz: A Global View of Transformer Attention*. ArXiv.org. <https://arxiv.org/abs/2305.03210>
- [19] Kieval, P.H. Mapping representational mechanisms with deep neural networks. *Synthese* 200, 196 (2022). <https://doi.org/10.1007/s11229-022-03694-y>
- [20] Schrimpf, M., Kubilius, J., Hong, H., Majaj, N. J., Rajalingham, R., Issa, E. B., Kar, K., Bashivan, P., Prescott-Roy, J., Geiger, F., Schmidt, K., Yamins, D. L. K., & DiCarlo, J. J. (2018). *Brain-Score: Which artificial neural network for object recognition is most brain-like?* bioRxiv. <https://doi.org/10.1101/407007>
- [21] Lake, B. M., Salakhutdinov, R., & Tenenbaum, J. B. (2023). Human-like systematic generalization through a meta-learning neural network. *Nature*, 623, 115–121. <https://doi.org/10.1038/s41586-023-06668-3>
- [22] Grabinski, J., Keuper, J. & Keuper, M. Aliasing and adversarial robust generalization of CNNs. *Mach Learn* 111, 3925–3951 (2022). <https://doi.org/10.1007/s10994-022-06222-8>